

Introduction to C++ - lectures 1-2

Martin Robinson

2019

A few introductory remarks

- C++ is a common programming language for high performance scientific computing applications, but it is not the only choice. You should know what languages/platforms are available to you.
- It is generally useful to know at least one low-level language (e.g. C++) and one high-level language (e.g. Python, Matlab) for your work.
- C++ is useful for:
 - High performance applications. Concurrency, parallelism and compilation to efficient, optimised binaries is easier in C++
 - Software libraries. C++ is designed for zero-cost abstraction, so library authors can provide their functionality with minimal overhead.
 - Access to a vast array of third-party scientific libraries

- There are lots of resources for C++, particularly online.
 - Many books, e.g. Bjarne Stroustrup, “Programming: Principles and Practice Using C++”. Make sure to use editions published well after 2011.
 - Reference material: <http://www.cplusplus.com> and <http://cppreference.com>
 - Q&A sites, <http://stackoverflow.com> etc.
- Integrated Development Environments (IDEs) exist—e.g. CLion or Visual Studio. These are useful for beginners as well as managing large code development. A command-line editor such as `vim` can also be highly customised for C++ development

Programming paradigms in C++

C++ is a multi-paradigm language with the goal of allowing powerful abstractions with *zero* run-time cost

- *Statically typed*: Variables have a constant type determined at compile-time, rather than determined on-the-fly at run-type (e.g. MATLAB)
- *Procedural*: The C-style basis of C++, algorithm consists of procedures (i.e. functions or subroutines) containing a series of computational steps
- *Object-orientated*: break computational task into objects (i.e. classes) that expose behaviours (methods) and data (members) through a public interface

Programming paradigms in C++

Recent updates (e.g. C++11, 14, 17) focus on generic and functional paradigms

- *Generic*: algorithms are written with generic types or variables to be determined later (at compile-time). This is achieved in C++ through *templates*
- *Functional*: Not a *pure* functional language, but functional elements exist in C++, for example the Standard Template Library (STL) and lambda functions

Lecture 1 — The basics

General structure of a basic C++ program

```
#include <header1>
#include <header2>
int main() {
    line of code;
    // this is a comment, ignored by the compiler
    more code; // this is a comment as well
    /*
        * Multi-line
        * comment
    */
    return 0;
}
```

Things to note

- Header files are listed first. These are files that contain the functions needed for operations such as input, output and mathematical calculations
- There is a section of code that starts `int main()`
 - or `int main(int argc, char* argv[])` when command-line input is important
- This section of code is followed by more code enclosed between curly brackets, { and }
- Comments may easily be inserted into the code
- Lines of code that “do something” end with a semicolon ;
- Just before the closing curly bracket at the end of the code is a statement `return 0;`

A first C++ program

```
#include <iostream>
int main()
{
    std::cout << "Hello World\n";
    return 0;
}
```

[< compiler explorer >]

This program prints the text “Hello World” to the screen

- `iostream` is a header file that is needed when using input and output
- `std::cout` is a command that sends output to the console, i.e. the screen
- `'\n'` is a formatting command that starts a new line
- All statements (lines of the program) inside the curly brackets end with a semicolon
;

Compiling the code

- A key difference between Matlab and C++ is that before the code can be executed it must be compiled
- When using software such as the Compiler Explorer or CLion, this is often done automatically
- When using the Gnu compiler of Unix/Cygwin this code can be compiled by saving the code and typing

```
g++ -o hello_world hello_world.cpp
```

followed by return.

- This produces an executable called `hello_world` that can be executed by typing

```
./hello_world
```

- Make is a general-purpose Unix tool that generates *targets* from their *dependencies* according to a set of *rules*, or *commands*, that you supply. These targets, dependencies and rules are all specified in a file called Makefile. For example:

```
hello_world: hello_world.cpp
    g++ -o hello_world hello_world.cpp
```

- Type `make` in a directory containing a Makefile, which by default generates the first target in the Makefile.

- CMake is a higher-level build system, which takes project definitions written in its own syntax, and generates Makefiles (or other, similar tools) that will compile your project, for example:

```
cmake_minimum_required(VERSION 3.10)
project(hello_world)
set(CMAKE_CXX_STANDARD 14)
add_executable(hello_world hello_world.cpp)
```

- Run cmake to generate the Makefile, then make to do the compilation

```
$ cmake <path-to-CMakeLists.txt>
$ make
```

- CLion is a C++ Interactive Development Environment (IDE) that, among other things, can handle compilation of (simple) projects automatically
- It uses CMake to do this, creating a `CMakeLists.txt` within each project
- You can edit the `CMakeLists.txt` directly to specify custom build behaviour
- CLion will be the recommended IDE to use during this course, although you are free to use whatever tool you are most familiar with.

Numerical variables

Before a variable is used the type of variable must be declared. For example if the variables `i` and `j` are integers and `a` is a double precision floating point number the statements

```
int i, j;  
double a;
```

must be included in the program before these variables are used.

It is advisable to use `double` rather than `float` in scientific computing applications.

Some names, such as `int`, `for`, `return` may not be used as variable names because they are used by the language.

- These words are known as reserved words or keywords

The following code adds two integers and print the answer to screen:

```
int integer1 = 5;  
int integer2 = 10;  
int answer = integer1 + integer2;  
  
std::cout << "The sum of " << integer1 << " and "  
          << integer2 << " is " << answer << '\n';
```

The details of the `std::cout` statement will be explained later

More on built-in types

A variable may be initialised when defining the variable type, for example

```
int i = 5;
```

The value of more than one variable may be assigned in each statement:

```
int i, j;  
i = j = 3;
```

This will cause confusion because it means

```
i = ( j = 3 );
```

and not

```
( i = j ) = 3;
```


When assigning values to floating point variables it is good programming practice to write numbers with decimal points, i.e.

```
double x = 5.0;  
float y = 7.0f;
```

rather than

```
double x = 5;  
float y = 7;
```

If a quantity is a constant throughout the program it may be declared as such

```
const double density = 45.621;
```

Fundamental types

- Boolean type
 - `true` or `false`
- Integer types
 - `signed` or `unsigned`
 - `short`, `long`, `long long`
 - e.g. `short int`; `int`; `long int`; `unsigned short`; `unsigned long long`
- Floating point types
 - `float`
 - `double`
 - `long double`

Representable values

Some variable types and ranges are given below. Note the these are operating system dependent:

Variable type	C++ name	Range
-----	-----	-----
integer	int	-2^{31} to $2^{31} - 1$
integer	long int	-2^{63} to $2^{63} - 1$
unsigned integer	unsigned int	0 to $2^{32} - 1$
floating point	float	-3.4×10^{38} to 3.4×10^{38}
floating point	double	-1.8×10^{308} to 1.8×10^{308}
floating point	long double	-1.2×10^{4932} to 1.2×10^{4932}

The values on your specific system can be found in `std::numeric_limits`:

[< compiler explorer >]

Mathematical operations

There is a shorthand for some mathematical operations

Longhand	Shorthand
-----	-----
<code>a = a + b;</code>	<code>a += b;</code>
<code>a = a - b;</code>	<code>a -= b;</code>
<code>a = a * b;</code>	<code>a *= b;</code>
<code>a = a / b;</code>	<code>a /= b;</code>
<code>a = a + 1;</code>	<code>a++;</code> if <code>a</code> is an integer
<code>a = a - 1;</code>	<code>a--;</code> if <code>a</code> is an integer

Some example lines of code

```
float a, b;  
double d, e;  
a = 3.0;  
b = ( a * std::pow(a, 3) ) / 2.0;  
d = 4.0;  
e = 2.0 * std::sqrt(d);
```

`std::pow(x,y)` gives the value of x^y . The `std::` indicates the standard namespace.

`std::sqrt(d)` gives the square root of the variable `d`

When using mathematical functions such as `std::pow`, you need to

```
#include <cmath>
```

Division of an integer by another integer will almost certainly cause problems

An example is given in the following piece of code

```
int i = 5, j = 2, k;  
k = i / j;  
std::cout << k << '\n';
```

The variable `k` is an integer and so cannot store the true value, 2.5

Instead, it will store the value 2

Suppose an integer is divided by a variable of type `double` - or vice versa - and that the result returned is stored in a variable of type `int`, as shown in the code below.

The variable `k` in this code is unable to store the mathematically correct answer

```
double i = 5.0;  
int j = 2, k;  
k = i / j;  
std::cout << k << '\n';
```

Only construct mathematical operations that are on elements of the same type

An integer can be converted, or cast, to a different data type – see the next slide

Variables can be converted from one type to another, for example:

```
double i = 5.0;
double k;
int j = 2;
k = i / static_cast<double>(j);
std::cout << k << '\n';
```

In this example, `static_cast<double>(j)` makes the variable `j` behave as if it were a `double` variable.

ASCII characters

ASCII characters are numbers, uppercase letters, lowercase letters and some other symbols

These characters may be represented using the data type `char`

```
#include <iostream>
int main() {
    char letter;
    letter = 'a'; // note the single quotation marks

    std::cout << "The character is " << letter << '\n';

    return 0;
}
```

Boolean variables

These variables take the values `true` or `false`, and are of use when using `if` conditionals and `while` loops

They are used as follows:

```
bool flag = true;
```

Strings

A character is one letter or number, a string is an ordered collection of characters

For example, "C++" is a string consisting of the ordered list of characters "C", "+", and "+"

To use strings in C++ requires an extra header, as shown below:

```
#include <iostream>
#include <string>

int main() {
    std::string city = "Oxford"; // note the std:: and the " marks
    std::cout << city << '\n';
}
```


The if statement

Suppose you want to execute some code only if the condition $p > q$ is met

This is achieved using the following code:

```
if (p > q)
{
    statement1;
    statement2;
}
```

Note the indentation. This makes it clear which statements are executed in the body of the **if** statement.

If only one statement is to be executed curly brackets aren't strictly necessary.

For example, the following code will execute statement1 if the condition $p > q$ is met

```
if (p > q)
    statement1;
```

but this is considered poor software engineering practice, and instead you should write this code as

```
if (p > q)
{
    statement1;
}
```

The use of curly brackets makes it clear which statement(s) are to be executed.

Third example – more than one condition

```
if (p == 0) {  
    statement1;  
} else if (p < 0) {  
    statement2;  
    statement3;  
} else {  
    // p > 0  
    statement4;  
}
```

Fourth example – nested `if` statements

```
if (p < q) {  
    if (x >= y) {  
        statement1;  
    }  
}
```

Fifth example – more than one condition

```
if (p < q || x < y) {  
    statement1;  
}
```

statement1 is executed only if one or both of $p < q$ and $x < y$ is true - i.e. `||` is the logical OR operator.

Relational and logical operators

relation	operator
-----	-----
equal to	== (note: not "=" which is assignment)
not equal to	!=
greater than	>
less than	<
greater than or equal to	>=
less than or equal to	<=

logical condition operator

AND &&

OR ||

NOT !

Boolean variables may be used in `if` statements as follows

```
bool flag1 = true, flag2 = false;
if (flag1)
{
    std::cout << "Does print something" << '\n';
}
if (flag2)
{
    std::cout << "Doesn't print anything" << '\n';
}
if (!flag2)
{
    std::cout << "Does print something" << '\n';
}
```

The while loop

```
while ( x < 100.0 && i < 10 ) {  
    x += x;  
    i++;  
}
```

The condition `x < 100.0 && i < 10` is tested only at the beginning of the statements in the loop, and not after every statement.

For example if the loop is entered when $x = 99.0$ and $i = 1$, the loop will be executed completely once.

Loop won't be entered when $x \geq 100$: `x` and `i` will be unchanged.

If you need a loop to execute at least once, with a test at the end, use:

- `do { ... } while (condition)`

The for loop

The following loop executes the statements inside the loop 10 times.

```
for (int i=0; i<10; i++)  
{  
    statement1;  
    statement2;  
}
```

Note that `i` can be previously declared, or declared in the loop statement.

The for loop

for loops can be nested and run over variable indices:

```
for (int i=0; i<5; i++)  
{  
    for (int j=0; j<10; j++)  
    {  
        std::cout << "i = " << i << " j = " << j << '\n';  
    }  
}
```

Random Number generation

You can generate random numbers using the standard library

```
#include <iostream>
#include <random>
int main() {

    std::default_random_engine generator;
    std::uniform_real_distribution<double> dist(0,1);

    double my_rand = dist(generator);
    return 0;
}
```

[< compiler explorer >]; [< cpp reference >]

Fixed sized arrays

If the size of the array is known in advance (i.e. at compile-time) then it is better to use the fixed size array `std::array<T,N>`, where T is the type that the array will hold (e.g. `int`, `double`), and N is the length of the array.

```
#include <array>
std::array<int, 4> x; // x holds 4 int
std::array<double, 5> y; // y holds 5 double

// a 5 by 5 array
std::array<std::array<double,5>,5> z;
```

In contrast to Matlab and FORTRAN the indices of an array of length `n` start at 0 and end at `n-1`

Elements of the array are accessed by placing the indices in separate square brackets, for example

```
x[0] = 1;  
z[1][2] = 3.0;
```

Arrays can be initialised when they are declared, for example

```
std::array<double,3> array1 = {0.0, 1.0, 2.0};  
std::array<std::array<double,3>,2> array2 =  
    {{{0.0, 1.0, 2.0},  
      {3.0, 4.0, 5.0}}};
```

Note that the values of arrays may only be set using the curly bracket notation when they are declared - for example the code

```
std::array<int,3> array = {0, 1, 2};
```

is correct, but the code

```
std::array<int,3> array;  
array = {0, 1, 2};
```

is not correct.

Console output

Console output may be achieved by using `std::cout`

We have already seen that the statement

```
std::cout << "Hello World\n";
```

prints the text “Hello World” to the screen, followed by a newline.

The statements

```
int x = 1, y = 2;  
std::cout << "x = " << x << " and y = " << y << '\n';
```

give the following output:

```
x = 1 and y = 2
```

Note that any spaces required in the output must be included within quotation marks.

Some useful formatting commands are:

Command	Symbol
-----	-----
newline	\n
tab	\t
'	\'
"	\"
(bell)	\a

Sometimes, for example if the computer is busy doing a large volume of computation, the program may not print the output to the screen immediately. If immediate output is desirable then use `std::cout.flush()` after any `std::cout` commands

```
std::cout << "Hello World\n";  
std::cout.flush();
```

or use `std::endl`, which combines a newline with a flush:

```
std::cout << "Hello World" << std::endl;
```

Keyboard input

Keyboard input for numbers and characters is achieved using `std::cin`

The following statements prompts someone to enter their PIN

```
int pin;  
std::cout << "Enter your PIN, then hit RETURN\n";  
std::cin >> pin;
```

`cin` may be used to ask for more than one input at a time

```
int accountno, pin;  
std::cout << "Enter your account number then hit RETURN,\n";  
std::cout << "and then your PIN followed by RETURN\n";  
std::cin >> accountno >> pin;
```

Keyboard input for strings is slightly different. An example is given below

```
#include <iostream>
#include <string>
int main()
{
    std::string name;
    std::cout << "Enter your name and then hit RETURN\n";
    std::getline(std::cin, name);
    std::cout << "Your name is " << name << '\n';
    return 0;
}
```

Writing to file

When writing to file, an additional header function `fstream` is needed.

```
#include <iostream>
#include <fstream>
#include <string>
```

The file `output.dat` may be opened using the statement

```
std::ofstream out("output.dat");
```

We can then write to this file in a similar manner as writing to the screen, with the exception that `cout` is replaced by `out`

There are a number of formatting options available. The following prints data in scientific format:

```
#include <iostream>
#include <fstream>
int main() {
    double x = -1.0, y = 45.3275893627129, z = 0.00000001;
    std::ofstream out("output.dat");
    out.setf(std::ios::scientific|std::ios::showpos);
    out << x << " " << y << " " << z << '\n';
    out.close();
    return 0;
}
```

[< compiler explorer >]; [< cpp reference >]

Reading from file

Suppose the file `numbers.dat` has 3 columns of numbers. This file can be opened using the following code

```
#include <fstream>

double x, y, z;
std::string line;
std::ifstream input("numbers.dat");
assert(input.is_open());
```

`input.is_open()` returns `true` if the file was successfully opened

We can read the file like so:

```
#include <sstream>

while(std::getline(input, line)) {
    std::istringstream s(line);
    s >> x >> y >> z;
}

input.close();
```

`std::getline` gets the next line of the file and returns `false` if we are at the end of the file.

`std::istringstream` converts the `std::string` `line` to a stream (like `std::cin`) which can be used to separate out the three columns.

Tip - Use of assert statements for debugging

- `assert` statements can be used in code to confirm something you expect to be true
- For example, you may wish to confirm that a number you are about to take the square root of is non-negative
- If the condition is not met, the code aborts giving an error message that explains what went wrong
- To use `assert` statements you must `#include <cassert>`
- An example of the use of `assert` statements is given on the next slide

```
#include <iostream>
#include <cassert>
#include <cmath>
int main()
{
    double a;
    std::cout << "Enter a non-negative number\n";
    std::cin >> a;
    assert(a >= 0.0);
    std::cout << "The square root of a is " << std::sqrt(a) << '\n';
    return 0;
}
```

If the code on the previous slide is compiled and run, and the user enters “-5” at the prompt, the code will be terminated and the following error message given:

```
a.out:: program.cpp:10: int main(): Assertion 'a >= 0.0' failed
```

This error message tells us to look at line 10 in `program.cpp`, where the assertion in quotes failed the test.